Prompt Documentation

1) **Classification prompt:**
You are an AI triage assistant for a B2B software company. Return only valid JSON. Classify the customer message using exactly one category from this list only: Bug Report, Feature Request, Billing Issue, Technical Question, Incident/Outage. Use these rules: login failures, errors, broken pages, and problems after updates are Bug Report. Requests for new functionality are Feature Request. Invoice or pricing problems are Billing Issue. Questions about setup, configuration, or capability are Technical Question. Service down, dashboard not loading, or multiple users affected are Incident/Outage. Priority must be one of: Low, Medium, High. Confidence score must be an integer from 0 to 100.

- Explanation: This classification prompt was structured to make the first decision point in the workflow accurate, consistent, and easy to use in automation. The system prompt clearly defines the model's role, limits the category labels to the five required by the assessment, and adds simple decision rules so the model can distinguish between similar request types such as bug reports, technical questions, and incidents. The user prompt then injects the source and raw message dynamically while forcing a strict JSON response with only category, priority, and confidence score, which reduces parsing errors and makes the output reliable for downstream routing in n8n. The main tradeoff was choosing a constrained and highly structured prompt over a more flexible one, which improves consistency but may reduce nuance in borderline cases. With more time, the prompt could be improved by adding a few representative examples for each class, refining the confidence-scoring instructions so the model avoids overusing very high scores, and adding extra validation logic for ambiguous requests.

2) **Enrichment prompt:**
You are an AI triage assistant for a B2B software company. Return only valid JSON. Extract structured enrichment data from the customer message. You must return: core_issue as one sentence, identifiers as an object, and urgency_signal as a short phrase. Capture account IDs, invoice numbers, error codes, times, auth providers, and any other useful identifiers if present.

- Explanation: This enrichment prompt was structured to turn unstructured customer messages into a clean intermediate record that the routing logic could use reliably downstream. The system prompt clearly defines the extraction task and limits the output to three specific elements: core issue, identifiers, and urgency signal—so the model focuses only on the information needed for triage rather than generating extra text. The user prompt reinforces this by forcing a strict JSON schema with named identifier fields such as account ID, invoice number, error code, mentioned time, and auth provider, which makes the output much easier to parse in n8n and later flatten into a Google Sheet. The main tradeoff was choosing a fixed structure instead of a more flexible extraction approach, which improves consistency but may miss uncommon identifiers that do not fit neatly into the predefined fields. With more time, this prompt could be improved by adding representative extraction examples, expanding the schema for more edge cases, and adding a validation or fallback step for messages where the model is unsure whether a value should be placed in a named field or in the other array.

3) **Generate Summary Prompt:**
   You are an AI triage assistant for a B2B software company. Return only valid JSON. Write exactly 2 professional sentences for the receiving team. The summary must be concise, clear, action-oriented, and each sentence must end with a period.

   - Explanation: This summary prompt was structured to generate a short, professional handoff note that a receiving team could read quickly and act on immediately. The system prompt forces a strict JSON-only response and explicitly requires exactly two sentences, which keeps the output consistent and easy to parse in n8n while also matching the assessment requirement for a short human-readable summary. The user prompt provides all the key upstream fields, classification, enrichment, routing, and escalation details, so the model can produce a concise summary that explains both the issue and the action taken. The main tradeoff was choosing a tightly constrained format over a more natural or flexible summary style, which improves reliability but can make the wording slightly repetitive across records. With more time, this prompt could be improved by adding a few examples of strong summaries, varying the wording based on queue or issue type, and adding a light quality check to ensure the summary always reflects the most important operational detail.